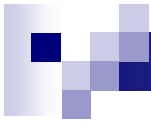




# Introduction to Parallel Computing

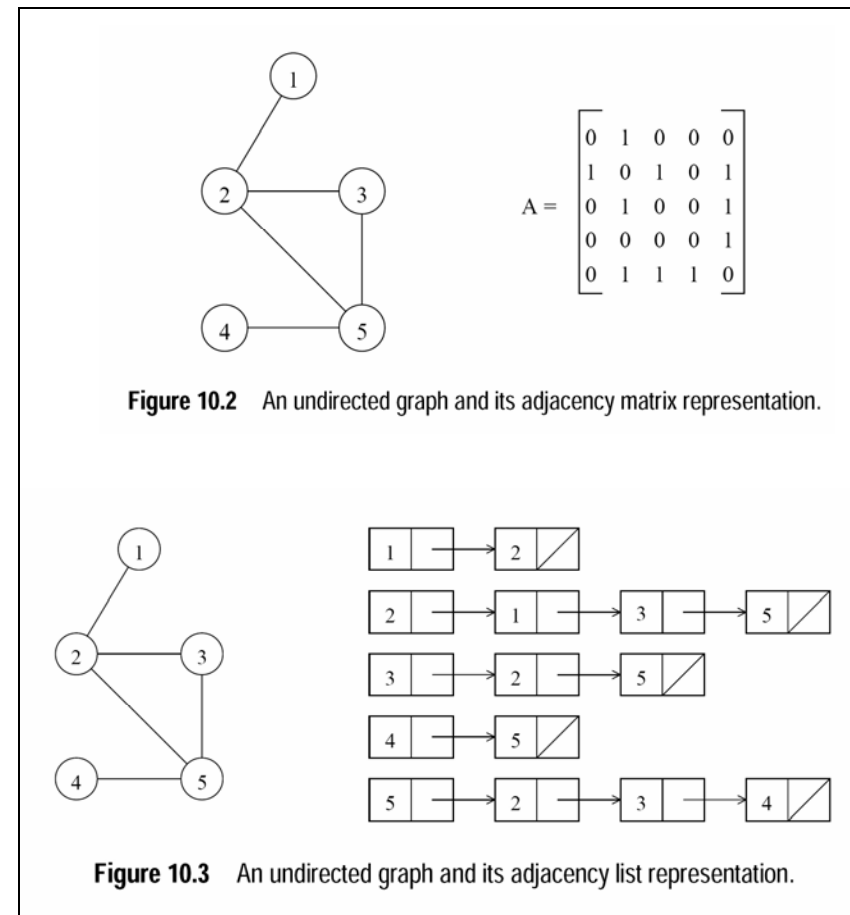

George Karypis  
Graph Algorithms



# Outline


- Graph Theory Background
- Minimum Spanning Tree
  - Prim's algorithm
- Single-Source Shortest Path
  - Dijkstra's algorithm
- All-Pairs Shortest Path
  - Dijkstra's algorithm
  - Floyd's algorithm
- Maximal Independent Set
  - Luby's algorithm

# Background



# Minimum Spanning Tree

- Compute the minimum weight spanning tree of an undirected graph.



# Prim's Algorithm


- Prim's Algorithm
  - $\Theta(n^2)$  serial complexity for dense graphs.
    - why?
- How can we parallelize this algorithm?
- Which steps can be done in parallel?

```

1.  procedure PRIM_MST( $V, E, w, r$ )
2.  begin
3.     $V_T := \{r\};$ 
4.     $d[r] := 0;$ 
5.    for all  $v \in (V - V_T)$  do
6.      if edge  $(r, v)$  exists set  $d[v] := w(r, v);$ 
7.      else set  $d[v] := \infty;$ 
8.    while  $V_T \neq V$  do
9.      begin
10.     find a vertex  $u$  such that  $d[u] := \min\{d[v] | v \in (V - V_T)\};$ 
11.      $V_T := V_T \cup \{u\};$ 
12.     for all  $v \in (V - V_T)$  do
13.        $d[v] := \min\{d[v], w(u, v)\};$ 
14.     endwhile
15.  end PRIM_MST

```

$d[]$  - for every node not in the tree what's its minimal distance to the tree



# Parallel Formulation of Prim's Algorithm

- Parallelize the inner-most loop of the algorithm.
  - Parallelize the selection of the “minimum weight edge” connecting an edge in  $V_T$  to a vertex in  $V - V_T$ .
  - Parallelize the updating of the  $d[]$  array.
- What is the maximum concurrency that such an approach can use?
- How do we “implement” it on a distributed-memory architecture?

```
1.  procedure PRIM_MST( $V, E, w, r$ )
2.  begin
3.     $V_T := \{r\};$ 
4.     $d[r] := 0;$ 
5.    for all  $v \in (V - V_T)$  do
6.      if edge  $(r, v)$  exists set  $d[v] := w(r, v);$ 
7.      else set  $d[v] := \infty;$ 
8.    while  $V_T \neq V$  do
9.      begin
10.        find a vertex  $u$  such that  $d[u] := \min\{d[v] | v \in (V - V_T)\};$ 
11.         $V_T := V_T \cup \{u\};$ 
12.        for all  $v \in (V - V_T)$  do
13.           $d[v] := \min\{d[v], w(u, v)\};$ 
14.        endwhile
15.      end PRIM_MST
```

# Parallel Formulation of Prim's Algorithm

- Decompose the graph  $A$  (adjacency matrix) and vector  $d$  vector using a 1D block partitioning along columns.
  - Why columns?
- Assign each block of size  $n/p$  to one of the processors.
- How will lines 10 & 12—13 be performed?
- Complexity? **Memory Complexity?**

highest parallelization possible?

What to do if  $\log\{p\} > n/p$  ?


“isoefficiency”

$$T_P = \overbrace{\Theta\left(\frac{n^2}{p}\right)}^{\text{computation}} + \overbrace{\Theta(n \log p)}^{\text{communication}}.$$

Find min  $d[]$  in  $O(\log p)$  communication steps  
 - on shared memory  
 - on a distributed architecture

```

1.  procedure PRIM_MST( $V, E, w, r$ )
2.  begin
3.     $V_T := \{r\};$ 
4.     $d[r] := 0;$ 
5.    for all  $v \in (V - V_T)$  do
6.      if edge  $(r, v)$  exists set  $d[v] := w(r, v);$ 
7.      else set  $d[v] := \infty;$ 
8.    while  $V_T \neq V$  do
9.      begin
10.     find a vertex  $u$  such that  $d[u] := \min\{d[v] | v \in (V - V_T)\};$ 
11.      $V_T := V_T \cup \{u\};$ 
12.     for all  $v \in (V - V_T)$  do
13.        $d[v] := \min\{d[v], w(u, v)\};$ 
14.     endwhile
15.  end PRIM_MST
    
```



# Single-Source Shortest Path

$l[i]$  - entry  $i$  is current distance of  $i$  to source  $s$  through tree nodes

- Given a *source* vertex  $s$  find the shortest-paths to all other vertices.
- Dijkstra's algorithm.
- How can it be parallelized for dense graphs?

---

```
1.  procedure DIJKSTRA_SINGLE_SOURCE_SP( $V, E, w, s$ )
2.  begin
3.     $V_T := \{s\};$ 
4.    for all  $v \in (V - V_T)$  do
5.      if  $(s, v)$  exists set  $l[v] := w(s, v);$ 
6.      else set  $l[v] := \infty;$ 
7.    while  $V_T \neq V$  do
8.      begin
9.        find a vertex  $u$  such that  $l[u] := \min\{l[v] | v \in (V - V_T)\};$ 
10.        $V_T := V_T \cup \{u\};$ 
11.       for all  $v \in (V - V_T)$  do
12.          $l[v] := \min\{l[v], l[u] + w(u, v)\};$ 
13.       endwhile
14.    end DIJKSTRA_SINGLE_SOURCE_SP
```

---

**Algorithm 10.2** Dijkstra's sequential single-source shortest paths algorithm.





# All-pairs Shortest Paths

- Compute the shortest paths between all pairs of vertices.
- Algorithms
  - Dijkstra's algorithm
    - Execute the single-source algorithm  $n$  times.
  - Floyd's algorithm
    - Based on dynamic programming.

# All-Pairs Shortest Path

## Dijkstra's Algorithm

### ■ Source-partitioned formulation

- Partition the sources along the different processors.

#### ■ Is it a good algorithm?

- Computational & memory scalability
- What is the maximum number of processors that it can use?

For  $p \leq n$ !  
Complexity  $T_p$ ?  
Memory?

### ■ Source-parallel formulation

- Used when  $p > n$ .
- Processors are partitioned into  $n$  groups each having  $p/n$  processors.
- Each group is responsible for one single-source SP computation.
- Complexity?

$$T_P = \overbrace{\Theta\left(\frac{n^3}{p}\right)}^{\text{computation}} + \overbrace{\Theta(n \log p)}^{\text{communication}}.$$



# Floyd's Algorithm

- Solves the problem using a dynamic programming algorithm.
  - Let  $d^{(k)}_{i,j}$  be the shortest path distance between vertices  $i$  and  $j$  that goes only through vertices  $1, \dots, k$ .

$$d^{(k)}_{i,j} = \begin{cases} w(v_i, v_j) & \text{if } k = 0 \\ \min \{ d^{(k-1)}_{i,j}, d^{(k-1)}_{i,k} + d^{(k-1)}_{k,j} \} & \text{if } k \geq 1 \end{cases}$$

```
1. procedure FLOYD_ALL_PAIRS_SP( $A$ )
2. begin
3.    $D^{(0)} = A$ ;
4.   for  $k := 1$  to  $n$  do
5.     for  $i := 1$  to  $n$  do
6.       for  $j := 1$  to  $n$  do
7.          $d^{(k)}_{i,j} := \min (d^{(k-1)}_{i,j}, d^{(k-1)}_{i,k} + d^{(k-1)}_{k,j})$ ;
8.   end FLOYD_ALL_PAIRS_SP
```

- Complexity:  $\Theta(n^3)$ .
  - Note: The algorithm can run in-place.
- How can we parallelize it?

# Parallel Formulation of Floyd's Algorithm

On k'th iteration matrix[i,j] is shortest distance from i to j going through nodes 1..k only

- Distribute the matrix using a 2D block decomposition.
- Parallelize the double inner-most loop.

```

1.  procedure FLOYD_ALL_PAIRS_SP(A)
2.  begin
3.       $D^{(0)} = A$ ;
4.      for  $k := 1$  to  $n$  do
5.          for  $i := 1$  to  $n$  do
6.              for  $j := 1$  to  $n$  do
7.                   $d_{i,j}^{(k)} := \min(d_{i,j}^{(k-1)}, d_{i,k}^{(k-1)} + d_{k,j}^{(k-1)})$ ;
8.      end FLOYD_ALL_PAIRS_SP
    
```


- Communication pattern?
- Complexity?

$$T_P = \overbrace{\Theta\left(\frac{n^3}{p}\right)}^{\text{computation}} + \overbrace{\Theta\left(\frac{n^2}{\sqrt{p}} \log p\right)}^{\text{communication}}.$$

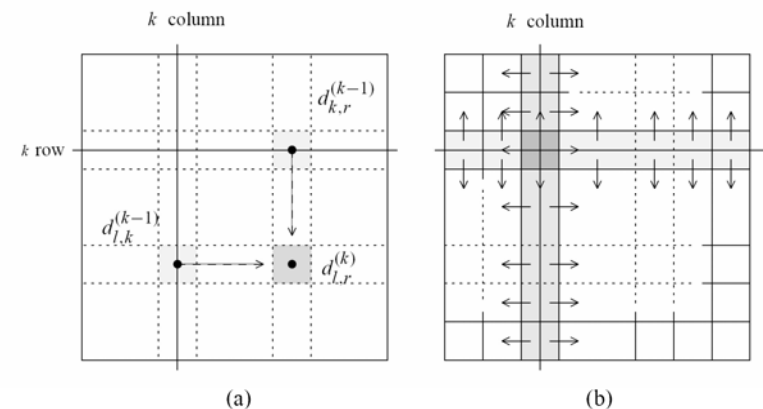
n iterations. each process performs per iteration:

communication - bcast  $n/\sqrt{p}$  values from k'th col/row in  $\log p$  steps

computation - modify  $n^2/p$  values



**Figure 10.7** (a) Matrix  $D^{(k)}$  distributed by 2-D block mapping into  $\sqrt{p} \times \sqrt{p}$  subblocks, and (b) the subblock of  $D^{(k)}$  assigned to process  $P_{i,j}$ .



**Figure 10.8** (a) Communication patterns used in the 2-D block mapping. When computing  $d_{i,j}^{(k)}$ , information must be sent to the highlighted process from two other processes along the same row and column. (b) The row and column of  $\sqrt{p}$  processes that contain the  $k$ 'th row and column send them along process columns and rows.

---

```
1.  procedure FLOYD_2DBLOCK( $D^{(0)}$ )
2.  begin
3.    for  $k := 1$  to  $n$  do
4.    begin
5.      each process  $P_{i,j}$  that has a segment of the  $k^{\text{th}}$  row of  $D^{(k-1)}$ ;
        broadcasts it to the  $P_{*,j}$  processes;
6.      each process  $P_{i,j}$  that has a segment of the  $k^{\text{th}}$  column of  $D^{(k-1)}$ ;
        broadcasts it to the  $P_{i,*}$  processes;
7.      each process waits to receive the needed segments;
8.      each process  $P_{i,j}$  computes its part of the  $D^{(k)}$  matrix;
9.    end
10. end FLOYD_2DBLOCK
```

---

**Algorithm 10.4** Floyd's parallel formulation using the 2-D block mapping.  $P_{*,j}$  denotes all the processes in the  $j^{\text{th}}$  column, and  $P_{i,*}$  denotes all the processes in the  $i^{\text{th}}$  row. The matrix  $D^{(0)}$  is the adjacency matrix.

# Comparison of All-Pairs SP Algorithms

**Table 10.1** The performance and scalability of the all-pairs shortest paths algorithms on various architectures with  $O(p)$  bisection bandwidth. Similar run times apply to all  $k-d$  cube architectures, provided that processes are properly mapped to the underlying processors.

	Maximum Number of Processes for $E = \Theta(1)$	Corresponding Parallel Run Time	Isoefficiency Function	what is n for computation time $\approx$ communication time
Dijkstra source-partitioned	$\Theta(n)$	$\Theta(n^2)$	$\Theta(p^3)$	
Dijkstra source-parallel	$\Theta(n^2 / \log n)$	$\Theta(n \log n)$	$\Theta((p \log p)^{1.5})$	
Floyd 1-D block	$\Theta(n / \log n)$	$\Theta(n^2 \log n)$	$\Theta((p \log p)^3)$	
Floyd 2-D block	$\Theta(n^2 / \log^2 n)$	$\Theta(n \log^2 n)$	$\Theta(p^{1.5} \log^3 p)$	
Floyd pipelined 2-D block	$\Theta(n^2)$	$\Theta(n)$	$\Theta(p^{1.5})$	


בהנחה שעלויות פעולות חישוב זהה לעלויות פעולות תקשורת, מהו מספר התהליכים הגדול ביותר כך שהגרעיניות (granularity) לא קטנה מידי (לא גורמת לעלויות תיאום גדולות מידי).

עבור מספר תהליכים כבעמודה השמאלית מהו זמן הריצה?

זהו מדד ליכולת ההרחבה של האלגוריתם: כאשר מגדילים את מספר התהליכים ( $p$ ) איך צריך להגדיל את בעיית החישוב ( $n$ ) כך שתישמר התאוצה (speedup).  
ככל ש-Isoefficiency קטנה יותר יכולת ההרחבה של האלגוריתם/מערכת (scalability) טובה יותר.

# Maximal Independent Sets

- Find the maximal set of vertices that are not adjacent to each other.



$\{a, d, i, h\}$  is <sup>Not</sup> an independent set

$\{a, c, j, f, g\}$  is a maximal independent set

$\{a, d, h, f\}$  is a maximal independent set

**Figure 10.15** Examples of independent and maximal independent sets.



# Serial Algorithms for MIS

- Practical MIS algorithms are incremental in nature.
  - Start with an empty set.
    1. Add the vertex with the smallest degree.
    2. Remove adjacent vertices
    3. Repeat 1—2 until the graph becomes empty.
- These algorithms are impossible to parallelize.
  - Why?
- Parallel MIS algorithms are based on the ideas initially introduced by Luby.



# Luby's MIS Algorithm

- Randomized algorithm.
  - Starts with an empty set.
  - 1. Assigns random numbers to each vertex.
  - 2. Vertices whose random number are smaller than all of the numbers assigned to their adjacent vertices are included in the MIS.
  - 3. Vertices adjacent to the newly inserted vertices are removed.
  - 4. Repeat steps 1—3 until the graph becomes empty.
- This algorithm will terminate in  $O(\log(n))$  iterations.
- Why is this a good algorithm to parallelize?
- How will the parallel formulation proceed?
  - Shared memory
  - Distributed memory

